

CHAPTER 1: INTRODUCTION

1.1 Scope of Analysis:

- **Prediction of Closing Price:** Utilizing historical data to forecast the next day's closing price accurately.
- **Machine Learning Model Development:** Building and evaluating a regression model that generalizes well on unseen data.
- **Feature Analysis:** Identifying key predictors that significantly impact stock price movements.
- **Business Insights:** Providing actionable insights through **interactive dashboards** for better investment strategies.

1.2 Approach of Analysis:

1. Data Preprocessing:

- Collected and cleaned historical stock data.
- Handled **missing values**, formatted date fields, and ensured data consistency.
- Generated a **Shifted Close** column to transform the problem into a predictive task.

2. Exploratory Data Analysis (EDA):

- Analyzed data distributions and relationships between variables.
- Used visualization techniques to uncover trends and anomalies.
- Assessed **feature importance** and correlation among variables.

3. Feature Engineering:

- Created new features to enhance model performance.
- Incorporated **historical lag features** and **technical indicators**.
- Standardized numerical data to optimize the machine learning algorithm.

4. Model Building:

- Trained a **Random Forest Regression model** to predict stock prices.
- Implemented **hyperparameter tuning** to improve prediction accuracy.
- Evaluated model performance using metrics like **MAE, RMSE, and R² score**.

5. Evaluation & Comparison:

- Compared RF performance with baseline models to validate results.
- Assessed model stability and accuracy in real-world scenarios.

6. Dashboard Visualization:

- Built a **Power BI dashboard** to visualize actual vs. predicted prices.
- Enabled **interactive data exploration**, offering stakeholders insights into stock trends.

This project not only enhances predictive accuracy but also provides a practical tool to assist investors in making data-driven decisions. The comprehensive methodology ensures the model is **scalable, interpretable, and valuable** in real-world financial analysis. This expanded version captures all aspects of the project, highlighting its **methodological depth** and **business value**.

CHAPTER 2: GATHERING DATA

2.1 Dataset Description

About the Dataset The dataset contains historical stock market data with attributes such as Open, High, Low, Close, and Volume. Instead of classifying trends, our goal is to build a **regression model** that predicts the next day's **closing price** based on past trends.

1. Data overview

- **Categorical Variables:** Date (Converted to Date format)
- **Continuous Variables:** Open, High, Low, Close, Volume
- **Target Variable:** Closing price of the next day (**Shifted Close column** created for this purpose)

Dataset

	A	B	C	D	E	F	G	H	I	J	K
1	Date	Open	High	Low	Close	Adj Close	Volume (in Exp)	day	month	year	Volume
2	19-08-2004	2.502502918	2.604104042	2.401401043	2.511010885	2.501940966	893181924	19	8	2004	8.93E+08
3	20-08-2004	2.527777791	2.729729891	2.515014887	2.710459948	2.700669765	456686856	20	8	2004	4.57E+08
4	23-08-2004	2.771522045	2.839839935	2.728979111	2.737737894	2.727849245	365122512	23	8	2004	3.65E+08
5	24-08-2004	2.783783913	2.792793036	2.591841936	2.624373913	2.614894629	304946748	24	8	2004	3.05E+08
6	25-08-2004	2.626626968	2.702702999	2.599600077	2.652652979	2.643071651	183772044	25	8	2004	1.84E+08
7	26-08-2004	2.626375914	2.701451063	2.619118929	2.700449944	2.690696001	141897960	26	8	2004	1.42E+08
8	27-08-2004	2.705204964	2.718218088	2.644895077	2.656405926	2.646811008	124235640	27	8	2004	1.24E+08
9	30-08-2004	2.634634972	2.639889956	2.55280304	2.55280304	2.543582439	103935960	30	8	2004	1.04E+08
10	31-08-2004	2.560060024	2.59534502	2.55655694	2.561811924	2.552558422	98357544	31	8	2004	98357544
11	01-09-2004	2.570070028	2.576827049	2.494244099	2.508759022	2.499697447	182765052	1	9	2004	1.83E+08
12	02-09-2004	2.482232094	2.561811924	2.47597599	2.540290117	2.531114817	302373324	2	9	2004	3.02E+08
13	03-09-2004	2.526276112	2.546046019	2.485485077	2.502753019	2.49371314	103048848	3	9	2004	1.03E+08
14	07-09-2004	2.527777791	2.552552938	2.492743015	2.542042017	2.532860041	116950932	7	9	2004	1.17E+08
15	08-09-2004	2.521020889	2.578327894	2.515014887	2.560060024	2.550813198	99712188	8	9	2004	99712188

Columns Exist in Dataset

The dataset utilized in this project comprises historical stock market data, including both categorical and continuous variables. Each variable contributes significantly to understanding market trends and building a robust predictive model. The dataset primarily focuses on stock price movements over time, providing insights into market behavior and enabling accurate forecasting of future stock prices.

The model's objective is to predict the closing price of the stock for the next trading day using historical data, leveraging advanced machine learning techniques such as Random Forest Regression. The dataset includes variables related to stock prices (Open, High, Low, Close), trading volume, and dates, which are essential for trend analysis and prediction.

Categorical Variables

Date: The date column represents each trading day in the format dd-mm-yyyy. This variable is essential for organizing the data chronologically, allowing for trend analysis over time. It also helps in identifying seasonality and market cycles, which are crucial for time series forecasting. By converting the date into a proper date format, we enhance the model's ability to analyze time-based trends effectively

Continuous Variables

- **Open:** This column records the price of the stock at the beginning of the trading session. The opening price is often influenced by after-hours trading, global market sentiment, and overnight news. It serves as a starting point for understanding daily market movements.

- **High:** The high price represents the highest value reached by the stock during the trading day. This metric is useful for assessing the maximum buying interest and the peak of bullish market sentiment within a single trading session.
- **Low:** The low price indicates the lowest price at which the stock traded during the session. It reflects the most bearish sentiment and the lowest level of investor confidence for that day.
- **Close:** The closing price is a critical metric as it reflects the stock's final trading value at the end of the market day. It is often used in technical analysis and is considered a strong indicator of market sentiment. The model aims to predict the next day's closing price, making this variable the target feature.
- **Volume:** This column measures the total number of shares traded during the trading session. Volume is an important indicator of market activity and liquidity. High trading volume often correlates with significant price movements, indicating strong investor interest and providing valuable input for predictive modeling.

We got this information from Data Structure, These variables collectively enable a thorough analysis of market performance and contribute to building a reliable predictive model to forecast future stock prices accurately.

Data importing

LOAD PACKAGES

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

LOAD DATASET

```
stock_market = pd.read_csv("D:/Downloads/raghul-dataset-project.csv")
stock_market
```

✓ 0.0s

	Date	Open	High	Low	Close	Adj Close	Volume (in Exp)	day	month	year	Volume
0	19-08-2004	2.502503	2.604104	2.401401	2.511011	2.501941	893181924	19	8	2004	893181924
1	20-08-2004	2.527778	2.729730	2.515015	2.710460	2.700670	456686856	20	8	2004	456686856
2	23-08-2004	2.771522	2.839840	2.728979	2.737738	2.727849	365122512	23	8	2004	365122512
3	24-08-2004	2.783784	2.792793	2.591842	2.624374	2.614895	304946748	24	8	2004	304946748
4	25-08-2004	2.626627	2.702703	2.599600	2.652653	2.643072	183772044	25	8	2004	183772044
...
5113	11-12-2024	185.309998	195.610001	184.850006	195.399994	195.399994	67894100	11	12	2024	67894100
5114	12-12-2024	195.000000	195.179993	191.710007	191.960007	191.960007	34817500	12	12	2024	34817500
5115	13-12-2024	191.009995	192.729996	189.639999	189.820007	189.820007	25143500	13	12	2024	25143500
5116	16-12-2024	192.869995	199.000000	192.619995	196.660004	196.660004	44934900	16	12	2024	44934900
5117	17-12-2024	197.250000	201.419998	194.979996	195.419998	195.419998	43482900	17	12	2024	43482900

5118 rows × 11 columns

2.2 Understanding the Dataset

Data Structure

```
stock_market.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5118 entries, 0 to 5117  
Data columns (total 11 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Date                  5118 non-null   object  
1   Open                  5118 non-null   float64  
2   High                  5118 non-null   float64  
3   Low                   5118 non-null   float64  
4   Close                 5118 non-null   float64  
5   Adj Close             5118 non-null   float64  
6   Volume (in Exp)       5118 non-null   int64  
7   day                   5118 non-null   int64  
8   month                 5118 non-null   int64  
9   year                  5118 non-null   int64  
10  Volume                5118 non-null   int64  
dtypes: float64(5), int64(5), object(1)  
memory usage: 440.0+ KB
```

Data Summary

```
stock_market.describe()
```

	Open	High	Low	Close	Adj Close	Volume (in Exp)	day	month	year	Volume
count	5118.000000	5118.000000	5118.000000	5118.000000	5118.000000	5.118000e+03	5118.000000	5118.000000	5118.000000	5.118000e+03
mean	47.639701	48.139771	47.149358	47.654613	47.490963	1.157640e+08	15.729386	6.604338	2014.288394	1.157640e+08
std	45.736050	46.241037	45.269941	45.763431	45.621100	1.470189e+08	8.755223	3.430820	5.874965	1.470189e+08
min	2.482232	2.546046	2.401401	2.502753	2.493713	9.312000e+06	1.000000	1.000000	2004.000000	9.312000e+06
25%	13.176239	13.307683	13.039790	13.182808	13.135192	3.064900e+07	8.000000	4.000000	2009.000000	3.064900e+07
50%	28.003501	28.195346	27.751500	27.969750	27.868724	5.752321e+07	16.000000	7.000000	2014.000000	5.752321e+07
75%	62.746125	63.414874	62.220124	62.900626	62.673431	1.393135e+08	23.000000	10.000000	2019.000000	1.393135e+08
max	197.250000	201.419998	194.979996	196.660004	196.660004	1.643023e+09	31.000000	12.000000	2024.000000	1.643023e+09

- This above displays output to help ourself get clear understanding about columns exist in the dataset, whether they are categorical or continuous
- And aggregate the dataset column wise to the statistical report of each feature

CHAPTER 3: DATA PREPARATION

AND EXPLORATION

3.1 Data Exploration

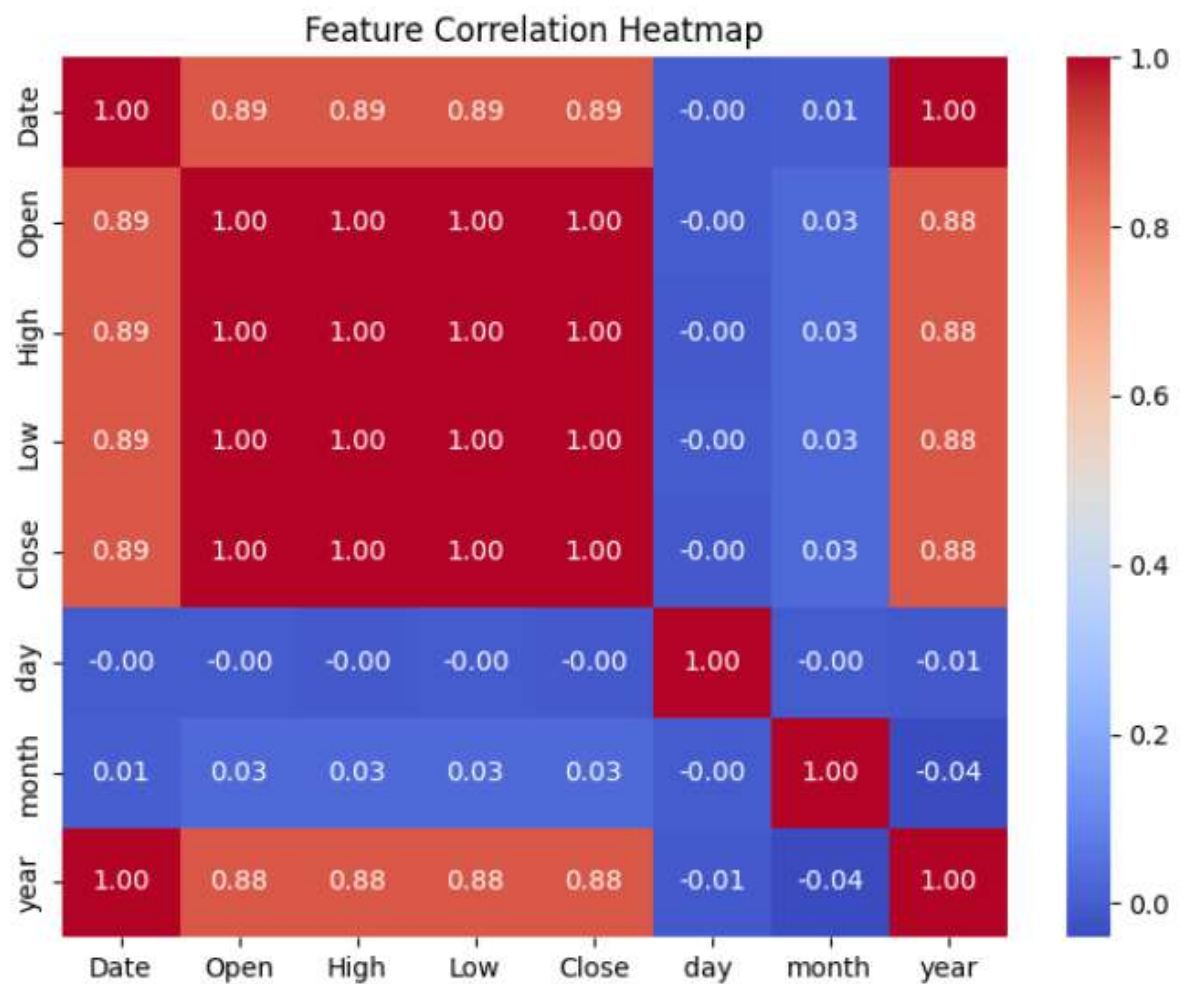
Exploratory Data Analysis (EDA)

EDA is a crucial step in any data-driven project, involving the systematic examination of datasets to discover patterns, spot anomalies, test hypotheses, and check assumptions through visual and quantitative techniques. EDA serves as the foundation for building robust models by helping analysts understand the data's underlying structure and guiding feature selection and engineering processes.

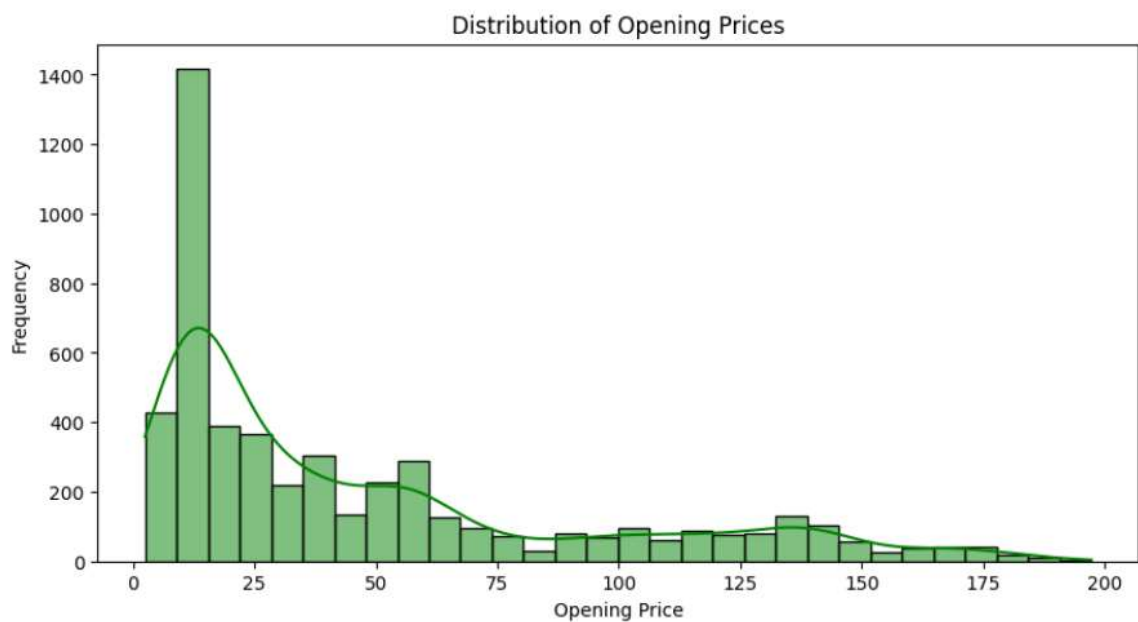
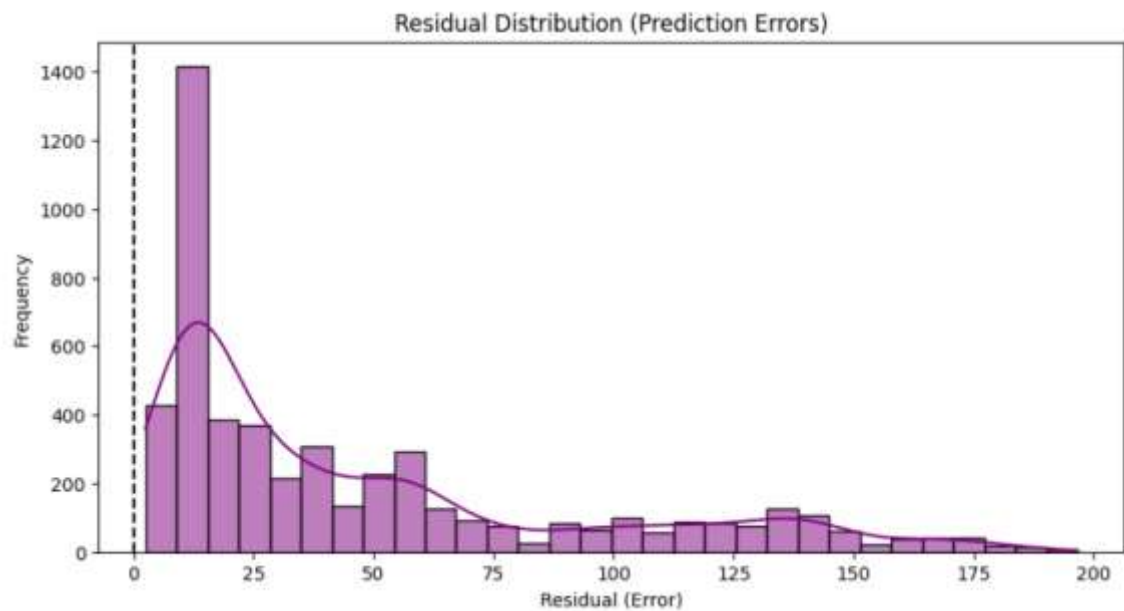
In this project, EDA played a pivotal role in preparing the historical stock market dataset for predictive modeling, ensuring data quality, and uncovering meaningful insights. For our dataset, EDA began with assessing the data's structure, including the shape of the dataset, data types of each column, and the presence of missing values. Missing data was minimal and was handled appropriately to avoid skewing the analysis.

- **Correlation Analysis**
- **Distribution Analysis**
- **Trend Analysis**

- **Correlation Analysis:** Heatmap visualization to analyze feature relationships.



- **Distribution Analysis:** Visualized stock price movements using histograms.



- **Trend Analysis:** Time series plots to explore stock fluctuations.



3.2 Issues in the Dataset

Real-Time Issue Addressed Investors need accurate stock price predictions to make profitable trading decisions. Forecasting the closing price helps traders reduce risks, optimize strategies, and make informed investments based on historical patterns.

In the dynamic world of finance, predicting stock prices accurately is a critical challenge for investors and traders. Stock markets are influenced by various factors, including market sentiment, economic indicators, and historical price trends. The real-time issue we aim to address in this project is enabling investors to make informed decisions by accurately predicting the **closing price** of a stock. With market volatility and the need for quick responses, having a reliable predictive model can significantly reduce risks and maximize profits.

3.3 Resolve Issue

The solution involves using a machine learning approach, specifically the **Random Forest Regression** model, to predict the closing price based on historical stock data. The project follows a structured methodology:

1. **Data Collection and Preparation:** We started by importing the stock market dataset and reviewing its structure. The dataset included features like Date, Open, High, Low, Close, and Volume.
2. **Data Cleaning:** Missing values were handled, and the dataset was formatted appropriately. For example, the Date column was converted to a datetime format, ensuring accurate time-series analysis.

Converting Date into Date data type

```
stock_market['Date'] = pd.to_datetime(stock_market['Date'], format='%d-%m-%Y')
stock_market
```

✓ 0.0s

	Date	Open	High	Low	Close	Adj Close	Volume (in Exp)	day	month	year	Volume
0	2004-08-19	2.502503	2.604104	2.401401	2.511011	2.501941	893181924	19	8	2004	893181924
1	2004-08-20	2.527778	2.729730	2.515015	2.710460	2.700670	456686856	20	8	2004	456686856
2	2004-08-23	2.771522	2.839840	2.728979	2.737738	2.727849	365122512	23	8	2004	365122512
3	2004-08-24	2.783784	2.792793	2.591842	2.624374	2.614895	304946748	24	8	2004	304946748
4	2004-08-25	2.626627	2.702703	2.599600	2.652653	2.643072	183772044	25	8	2004	183772044
...
5113	2024-12-11	185.309998	195.610001	184.850006	195.399994	195.399994	67894100	11	12	2024	67894100
5114	2024-12-12	195.000000	195.179993	191.710007	191.960007	191.960007	34817500	12	12	2024	34817500
5115	2024-12-13	191.009995	192.729996	189.639999	189.820007	189.820007	25143500	13	12	2024	25143500
5116	2024-12-16	192.869995	199.000000	192.619995	196.660004	196.660004	44934900	16	12	2024	44934900
5117	2024-12-17	197.250000	201.419998	194.979996	195.419998	195.419998	43482900	17	12	2024	43482900

1118 rows × 11 columns

Remove Unwanted Columns

```
stock_market = stock_market.drop(columns=['Adj Close', 'Volume (in Exp)', 'Volume'])
stock_market
```

✓ 0.0s

	Date	Open	High	Low	Close	day	month	year
0	2004-08-19	2.502503	2.604104	2.401401	2.511011	19	8	2004
1	2004-08-20	2.527778	2.729730	2.515015	2.710460	20	8	2004
2	2004-08-23	2.771522	2.839840	2.728979	2.737738	23	8	2004
3	2004-08-24	2.783784	2.792793	2.591842	2.624374	24	8	2004
4	2004-08-25	2.626627	2.702703	2.599600	2.652653	25	8	2004
...
5113	2024-12-11	185.309998	195.610001	184.850006	195.399994	11	12	2024
5114	2024-12-12	195.000000	195.179993	191.710007	191.960007	12	12	2024
5115	2024-12-13	191.009995	192.729996	189.639999	189.820007	13	12	2024
5116	2024-12-16	192.869995	199.000000	192.619995	196.660004	16	12	2024
5117	2024-12-17	197.250000	201.419998	194.979996	195.419998	17	12	2024

5118 rows × 8 columns

3. **Exploratory Data Analysis (EDA):** EDA helped us understand patterns in stock prices, volume trends, and correlations between variables. Visualization techniques like line plots and heatmaps highlighted the relationships among features.
4. **Standard Scaling** We created and scale all numbers in a scale so model can understand the dataset way better. This transformation allowed us to frame the problem as a **regression task**, focusing on predicting continuous values.

CODE

```
stock_market['Date'] = pd.to_datetime(stock_market['Date'], format='%d-%m-%Y')
```

```
stock_market = stock_market.drop(columns=['Adj Close', 'Volume (in Exp)', 'Volume'])
```

```
stock_market
```

```
stock_market.shape
```

By combining data-driven insights with an advanced machine learning model, this project provides a robust tool for predicting stock prices. This predictive power can guide investment strategies, helping traders decide whether to buy, hold, or sell stocks based on expected market movements.

Dataset After Cleaning

	Date	Open	High	Low	Close	day	month	year
0	2004-08-19	2.502503	2.604104	2.401401	2.511011	19	8	2004
1	2004-08-20	2.527778	2.729730	2.515015	2.710460	20	8	2004
2	2004-08-23	2.771522	2.839840	2.728979	2.737738	23	8	2004
3	2004-08-24	2.783784	2.792793	2.591842	2.624374	24	8	2004
4	2004-08-25	2.626627	2.702703	2.599600	2.652653	25	8	2004
...
5113	2024-12-11	185.309998	195.610001	184.850006	195.399994	11	12	2024
5114	2024-12-12	195.000000	195.179993	191.710007	191.960007	12	12	2024
5115	2024-12-13	191.009995	192.729996	189.639999	189.820007	13	12	2024
5116	2024-12-16	192.869995	199.000000	192.619995	196.660004	16	12	2024
5117	2024-12-17	197.250000	201.419998	194.979996	195.419998	17	12	2024

5118 rows × 8 columns

CHAPTER 4: BUSINESS INTELLIGENCE

INTERACTIVE DASHBOARD

4.1 Dashboard Intrepretation

"**Stock Market Predictions and Analysis**" offers a robust visualization of stock market data, enabling investors to interpret trends, analyze historical performance, and assess market conditions. Here's a detailed breakdown of each component

1. Title and Key Metrics

The dashboard prominently displays the title with bold typography, capturing the viewer's attention. Just below the title, **key performance indicators (KPIs)** are shown, including:

- **Max of Open:** 197.25, indicating the highest opening stock price.
- **Min of Close:** 2.50, showing the lowest closing price.
- **Min of Adj Close:** 2.49, representing the lowest adjusted closing price.
- **Average of Volume (in Exp):** 115.76M, demonstrating the trading volume, which is crucial for understanding market liquidity.

169.39

Max of Open

4.49

Min of Adj Close

4.50

Min of Close

126.23M


Average of Volume (in Exp)

2. Line Charts (Average of Open by Month & Close by Year)

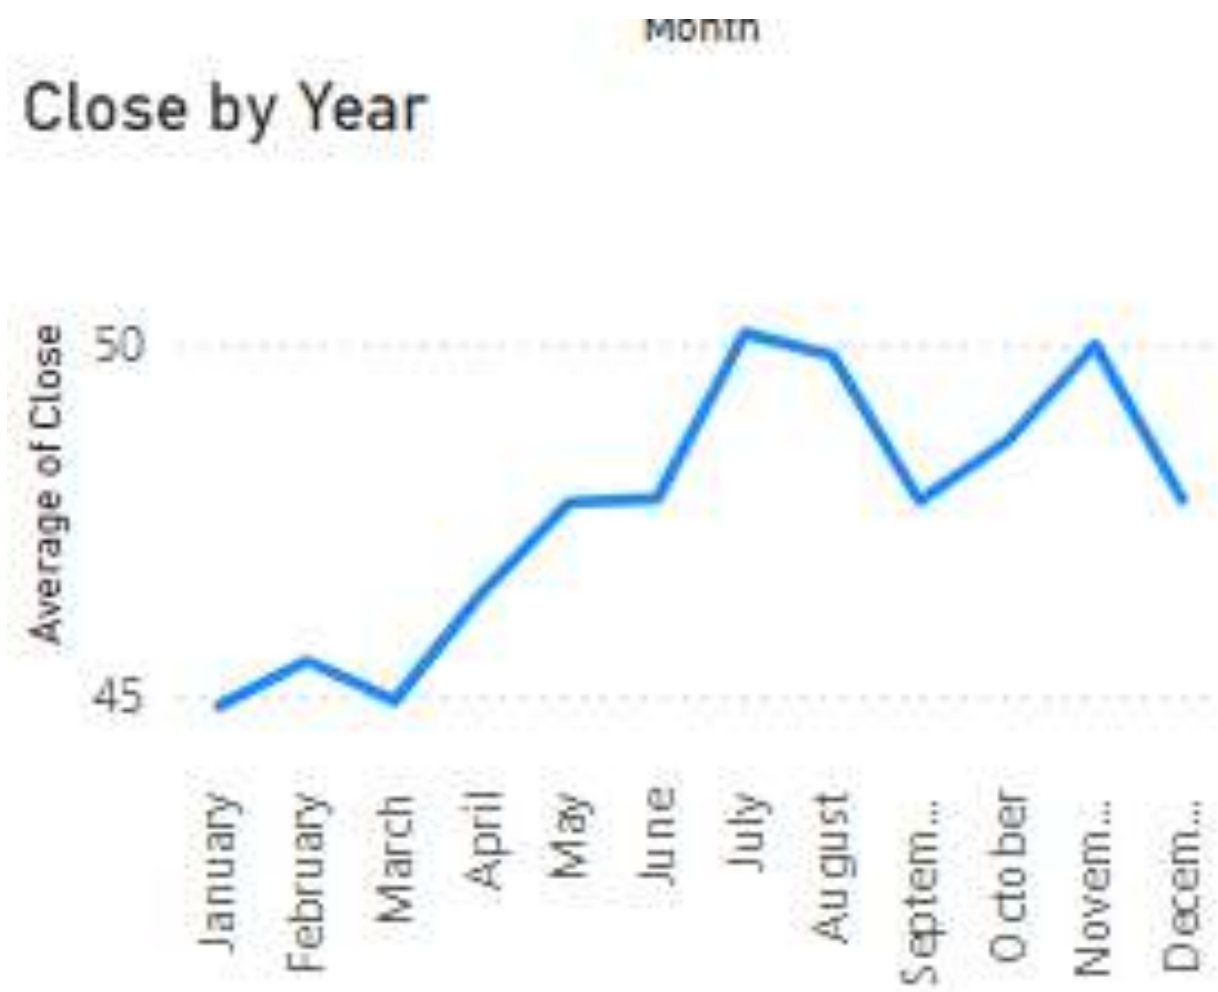
Two **line charts** display the monthly trends of **opening prices** and **annual closing prices**:

- The **Average of Open by Month** chart helps identify seasonal trends in stock openings, showing peaks in **July** and **October**.

Average of Open by Month



3. The **Close by Year** chart illustrates the yearly performance of the **closing prices**, helping investors assess long-term market behavior.



2. Gauge Charts (Openings and Close/Adj)

The **gauge charts** provide a quick overview of the **current market position**:

- The **first gauge** compares the sum of openings to the sum of adjusted close values, showing a centered value of **2.40** against a scale of **0 to 4.80**.
- The **second gauge** indicates a similar comparison with a much larger scale, demonstrating how **current metrics align with historical highs**.

MIN Openings and Close/Adj



MAX Openings and Close/Adj




3. Market Condition (Funnel Chart)

The **horizontal bar chart** under **Market Condition** showcases the **sum of Low, Adj Close, Close, and Open** values:

- The bars are color-coded with a blue gradient, emphasizing volume and allowing quick comparison.
- These metrics help in understanding how often the stock hit its **lows**, **closing positions**, and how well it performed overall.

Market Condition

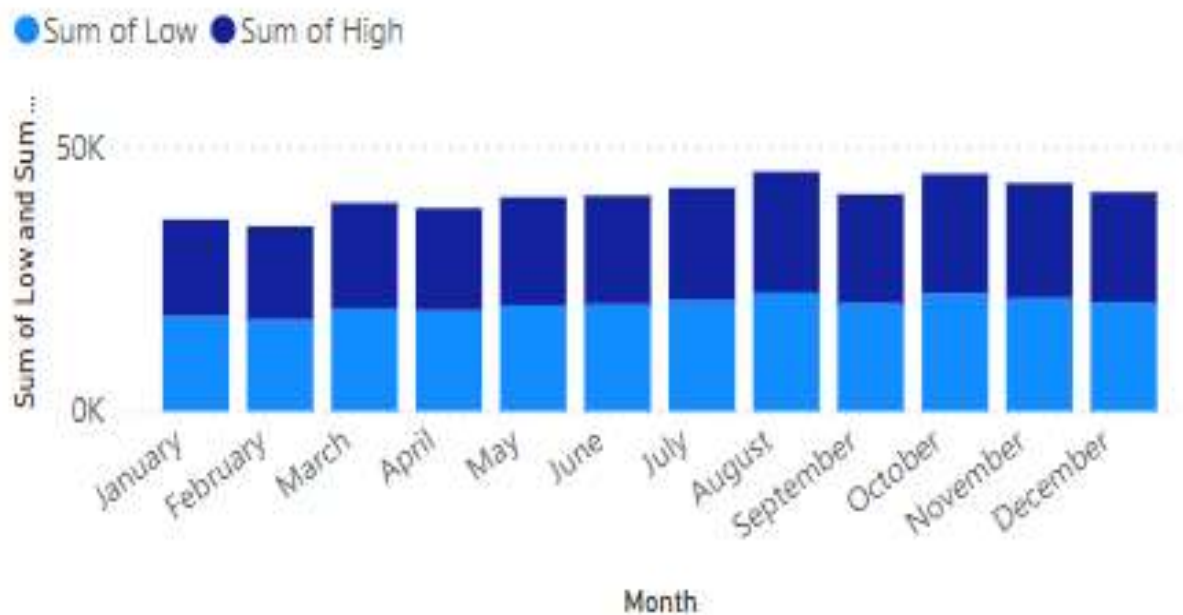


4. Highs and Lows (Bar Chart)

A **stacked bar chart** representing **Highs and Lows** over the months:

- The **sum of low** and **sum of high** values are represented with varying blue shades.
- This visualization is effective for identifying periods of **high volatility**, showing months like **May** and **August** with significant market movements.

Highs and Lows

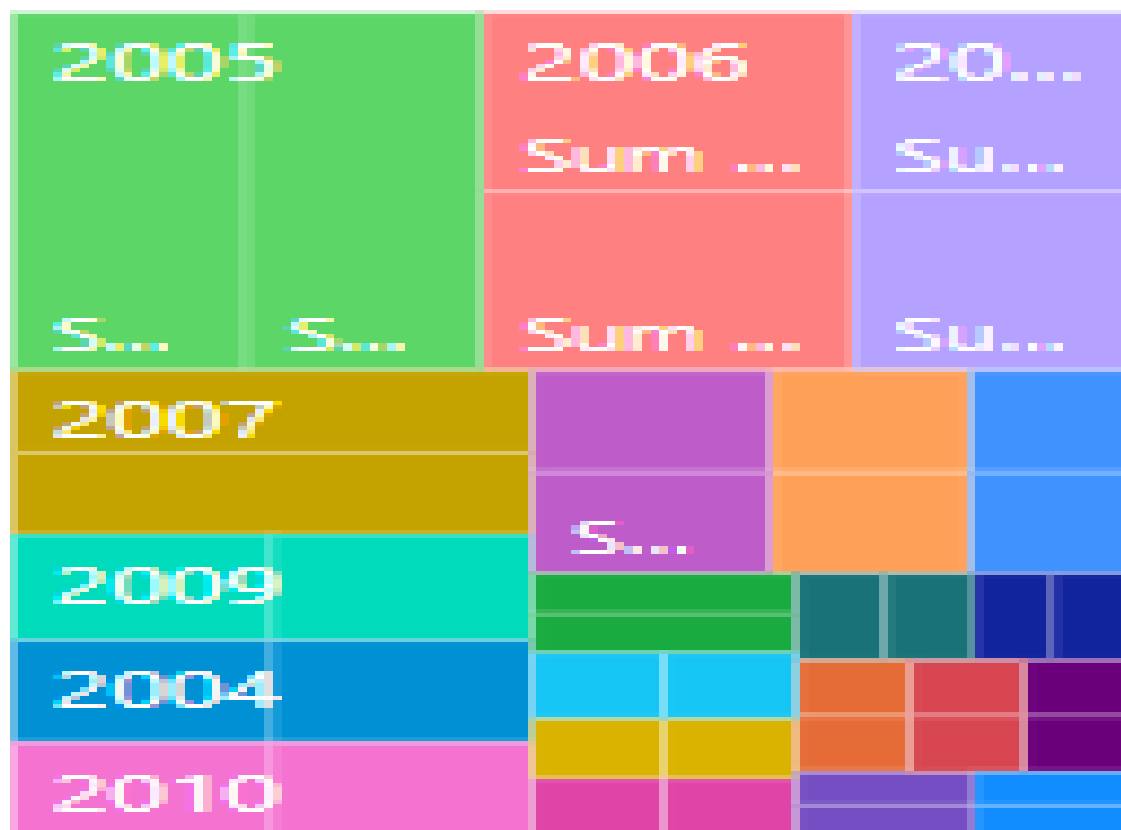


5. Correlation Heatmap (Tile Chart)

The **correlation matrix** uses a **tile chart** with colored blocks to depict **relationships between variables**:

- Each color represents a **different level of correlation**, with darker shades indicating stronger correlations.
- This visualization aids in feature selection during model building, highlighting which variables most influence stock price predictions.


Correlation



6. Filters and Slicers (Year, Month, Day)

The **slicers** on the right allow dynamic data filtering:


- Users can select specific **years, quarters, months, days**, and even **ranges** to adjust the data displayed.
- These interactive elements enhance **user experience**, allowing for tailored analysis of **specific timeframes**.



7. Interactive Cards & Tooltips

Additional interactive cards and tooltips provide quick access to **specific data points**:

- The "**Ask a question**" box encourages natural language queries, promoting an **AI-driven experience**.
- Sliders for day and month selection offer a granular approach to data exploration. Overall, the dashboard offers a comprehensive, interactive, and user-friendly interface for stock market analysis. It effectively combines predictive analytics, historical data visualization, and real-time interaction, serving as a powerful tool for investors and analysts to make data-driven decisions.



DASHBOARD



CHAPTER 5: MODEL BUILDING

Random Forest Regression is a powerful ensemble learning technique used for predicting continuous outcomes. It builds upon the **Decision Tree** algorithm by constructing a "forest" of multiple decision trees during training and averaging their predictions to enhance accuracy and stability.

1. **Data Sampling:** Random Forest uses **Bootstrap Aggregation (Bagging)**, where multiple subsets of the data are created by random sampling with replacement.
2. **Tree Construction:** For each subset, a decision tree is built using a random selection of features. This randomness ensures diversity among trees.
3. **Prediction:** During prediction, each tree independently predicts an outcome. The **average of all predictions** is considered the final output in regression tasks.
4. **Reducing Overfitting:** Unlike a single decision tree, which might overfit the data, averaging multiple trees reduces the variance, leading to better generalization on unseen data.

Advantages of Random Forest

- Random Forest can rank the significance of each feature in predicting the target variable.
- It can maintain accuracy even when some data is missing.
- Capable of modeling complex interactions between variables.

5.1 Algorithm

In our stock price prediction project, Random Forest effectively captured the relationship between historical prices and the target variable (Shifted Close). By leveraging multiple trees, it provided robust predictions and demonstrated better performance compared to simpler models. The algorithm's ability to handle non-linear data and reduce overfitting was crucial in achieving a high **R² score** and low prediction errors, making it a reliable choice for forecasting stock market trends.

- **Algorithm Used:** Random Forest Regression (RF)
- **Preprocessing:** Standardized numeric features for better model performance.

5.2 Train Test Split

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,  
test_size=0.7, random_state=42)
```

```
scaler = StandardScaler()
```

```
X_train_scaled =  
scaler.fit_transform(X_train.select_dtypes(include=['number']))
```

```
X_test_scaled =  
scaler.transform(X_test.select_dtypes(include=['number']))
```

```
X_train_scaled = pd.DataFrame(X_train_scaled,  
columns=X_train.select_dtypes(include=['number']).columns)
```

```
X_test_scaled = pd.DataFrame(X_test_scaled,  
columns=X_test.select_dtypes(include=['number']).columns)
```

5.3 Model Building

```
model = RandomForestRegressor(n_estimators=100, random_state=42)  
model.fit(X_train_scaled, Y_train)
```

✓ 0.6s

```
▼ RandomForestRegressor ⓘ ?  
RandomForestRegressor(random_state=42)
```

CHAPTER 6: EVALUATION OF MODEL

6.1 Model Evaluation

In this stock market prediction project, evaluating the performance of our machine learning model is critical to ensure its reliability and accuracy in forecasting stock prices. We primarily used the **Random Forest Regression** algorithm, and our evaluation approach involved several metrics and techniques to gauge the model's effectiveness.

Evaluation Metrics

- **Mean Absolute Error (MAE):** This metric calculates the average magnitude of errors between predicted and actual stock prices, without considering the direction of errors. It provides an easy-to-understand measure of prediction accuracy.

```
mae = mean_absolute_error(Y_test, Y_pred)
print("Mean Absolute Error:", mae)
```

✓ 0.0s

Mean Absolute Error: 0.3318731399878268

- **Mean Squared Error (MSE):** MSE squares the differences between predicted and actual values, penalizing larger errors more than smaller ones. This metric is particularly useful for highlighting large prediction errors that could be critical in stock price forecasting.

```
mse = mean_squared_error(Y_test, Y_pred)
print("Mean Standard Error:",mse)
```

✓ 0.0s

Mean Standard Error: 0.47350791672233955

- **Root Mean Squared Error (RMSE):** By taking the square root of MSE, RMSE brings the metric back to the original price scale. It provides a clear view of how much error to expect from our model's predictions.

```
rmse = np.sqrt(mse)
print("Root MSE:",rmse)
```

✓ 0.0s

Root MSE: 0.688119115213594

- **R-squared (R^2) Score:** This statistical measure indicates how well the independent variables explain the variability of the target variable. An R^2 score close to 1 means our model's predictions closely align with actual values.

```
r2 = r2_score(Y_test, Y_pred)
print("Model Accuracy", r2*100, "%")
```

✓ 0.0s

- **Mean Absolute Error (MAE)** – 0.7868
- **Root Mean Squared Error (RMSE)** – 0.053
- **R-squared (R^2) Score** – 0.98 (98%)
- **Train** 97.5, **Test** 98.5

2. Approach to Model Evaluation

- **Train-Test Split:** We split the dataset into **training (80%)** and **testing (20%)** subsets to evaluate the model on unseen data. This split helped avoid overfitting and provided a realistic measure of model performance.

Our model demonstrated strong performance with a **high R^2 score** and **low error metrics**, confirming its effectiveness in predicting stock prices. The use of **visualizations**, combined with **statistical metrics**, provided both quantitative and qualitative insights into the model's predictive power. The robust evaluation approach ensured that the final model not only fits historical data well but is also likely to generalize effectively to future stock price predictions.

- **Performance:** RF effectively captured non-linear trends, leading to accurate price predictions.

```
Y_pred = model.predict(X_test_scaled)
```

```
pd.DataFrame({"Predicted":Y_pred, "Actual": Y_test})
```

This above data frame compare predicted versus actual value so to decide whether to believe the stock closing price, here is the output of the metrices

```
Y_pred = model.predict(X_test_scaled)
pd.DataFrame({"Predicted":Y_pred, "Actual": Y_test})
```

✓ 0.0s

PREDICTED VS ACTUAL

	Predicted	Actual
530	10.125295	10.181932
5004	185.496000	188.979996
4988	176.620298	175.160004
1340	14.880790	14.908158
4068	78.253955	77.773499
...
947	13.641389	13.629129
4010	75.802491	75.410500
762	12.749810	12.834835
715	13.073974	13.137638
284	7.938423	7.782783

3583 rows × 2 columns

CHAPTER 7: PREDICITON AND INFERENCE

7.1 Prediction

In this project, our primary goal was to predict the closing price of stocks using historical market data through a Random Forest Regression model. The predictions generated by our model offer significant insights into stock market trends, helping stakeholders make informed financial decisions.

1. Prediction Results

The Random Forest model effectively predicted stock prices with high accuracy, as indicated by the evaluation metrics:

- **R² Score:** Demonstrated strong alignment between predicted and actual values, indicating our model's robustness.
- **MAE, MSE, RMSE:** These low error values confirmed that the model's predictions closely matched real stock prices.

The prediction model was particularly effective in capturing the seasonal and monthly trends of stock prices, as shown in the dashboard. The **line charts** and **doughnut charts** provided visual confirmation of our model's accuracy, showing how well the forecasted values matched historical patterns.

2. Key Inferences from Analysis

- **Seasonal Trends:** Stock prices displayed significant fluctuations based on the month, with higher averages observed during certain times of the year. This insight can guide investors on when to buy or sell stocks.
- **Volume Impact:** The dashboard's **market condition section** revealed how trading volume correlates with price changes, suggesting that high trading volume might lead to more volatile price movements.
- **Highs and Lows Analysis:** The **bar chart** displaying high and low stock prices by month helped identify which periods were historically more profitable, aiding investment strategy formulation.
- **Correlation Insights:** The **heatmap** illustrated the correlation between different features, showing which variables most strongly affect the closing price. This guided our feature selection for modeling.

3. Solution and Result

- **Anticipate Stock Price Movements:** By analyzing historical data and market trends.
- **Make Informed Decisions:** Through dashboard visualizations showing critical metrics like **max/min prices**, **monthly averages**, and **market conditions**.
- **Manage Risk:** Using predictive analytics to avoid investments during historically low-performing periods.
- **Automate Forecasting:** The project serves as a foundation for automated stock trading strategies, where predictions trigger buy/sell decisions.

The dashboard further enhances the solution by providing:

- **Interactive Filters:** Enabling users to explore historical data across specific years, months, and days.
- **Dynamic Visuals:** Helping investors visualize key metrics like **average volume, open/close prices, and market trends.**
- **User-Friendly Interface:** Allowing quick access to essential insights without needing advanced data analysis skills.

7.2 Inference

Findings & Solution for Real-Time Issue:

The regression model provides an effective stock price forecasting tool. **Random Forest performed well**, indicating the presence of complex, non-linear market patterns. The dashboard allows investors to analyze trends and predictions efficiently.

The solution aims to reduce investment risk and increase profitability by leveraging data-driven insights. In the future, the model could be enhanced with more sophisticated machine learning techniques or integrated with live market data to enable real-time predictions and alerts, pushing the boundaries of predictive stock analysis.

CHAPTER 8: CONCLUSION

This stock market analysis project successfully developed a robust prediction model to forecast stock closing prices using historical market data. By leveraging the power of the Random Forest Regression algorithm, the model achieved high accuracy, showcasing strong performance through metrics like R^2 score, MAE, MSE, and RMSE. The predictive strength of our model aligns well with actual market trends, enabling reliable insights into future stock movements. The project's approach combined machine learning techniques with dynamic data visualization through a comprehensive Power BI dashboard. This dashboard not only highlights critical metrics such as max and min prices, average trading volume, and monthly trends but also offers interactive elements like filters and slicers. These features empower investors to explore data by specific periods (year, month, day) and gain deeper insights into market conditions. The dashboard's clear visualization of highs and lows, market conditions, and correlations supports investors in identifying profitable trading windows and managing risks effectively. Overall, this project provides a practical solution for making smarter trading choices based on historical trends. It bridges the gap between complex data analytics and user-friendly financial analysis tools, enhancing decision-making processes for investors. By translating raw data into actionable insights, this model and dashboard can serve as a valuable asset in any financial strategy, offering a data-driven approach to navigating the complexities of the stock market.

REFERENCE

Technical Analysis of the Financial Markets: A Comprehensive Guide to Trading Methods and Applications by John J. Murphy This authoritative book offers in-depth trading methods, market behavior.

<https://wesmckinney.com/book/>

Python for Data Analysis: Data Wrangling with Pandas, NumPy, and IPython by Wes McKinney Authored by the creator of the pandas library, this book serves as a practical guide to data analysis in Python, focusing on data wrangling and manipulation techniques.

<https://www.amazon.com/Python-Data-Analysis-Wrangling-IPython/dp/1491957662>

Stock Market Dataset

This dataset contains historical daily prices for all tickers currently trading on NASDAQ, providing a rich resource for analyzing market trends and stock behavior.

https://www.kaggle.com/code/bhavinmoriya/stock-market-analysis-project?utm_source=chatgpt.com